
MAST ANALYSIS OF A RAVENSCAR APPLICATION WITH FPS AND EDF SCHEDULING

TECHNICAL REPORT

Giovanni Jiayi Hu

Department of Mathematics
University of Padua, Italy I-35121

Email: giovannijiayi.hu@studenti.unipd.it

Alessio Gobbo

Department of Mathematics
University of Padua, Italy I-35121

Email: alessio.gobbo@studenti.unipd.it

December 6, 2019

ABSTRACT

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

1 Introduction

Embedded systems have to satisfy strict timing requirements and especially in the case of such hard real-time applications, predictability of the timing behavior is an extremely important aspect.

The choice of a suitable design and development method, in conjunction with supporting tools that enable the real-time performance of a system to be analysed and simulated, can lead to a high level of confidence that the final system meets its real-time constraints.

As a matter of fact, the use of Ada has proven to be of great value within high integrity and real-time applications, thanks to language subsets of deterministic constructs, to ensure full analysability of the code. In the next sections we will use the term [RM] to refer to a section of the Ada Reference Manual¹.

Notably, the Ravenscar profile [1] is a subset of the tasking model, restricted to meet the real-time community requirements for determinism, schedulability analysis and memory-boundedness, as well as being suitable for mapping to a small and efficient run-time system that supports task synchronization and communication.

Along with the Ravenscar profile, we have used a model for representing the temporal and logical elements of real-time applications, called MAST [3]. This model allows a very rich description of the system, including the effects of event or message-based synchronization, multiprocessor and distributed architectures as well as shared resource synchronization.

1.1 The application

The example application presented in this paper is extracted from "Guide for the use of the Ada Ravenscar Profile in high integrity systems" [1]. It includes a periodic process that handles orders for a variable amount of workload. Whenever the request level exceeds a certain threshold, the periodic process farms the excess load out to a supporting

¹http://www.ada-auth.org/standards/rm12_w_tc1/html/RM-TOC.html




Figure 1: Architecture of the example application [1].

sporadic process. While such orders are executed, the system may receive interrupt requests from an external source. Each interrupt treatment records an entry in an activation log.

When specific conditions hold, the periodic process releases a further sporadic process to perform a check on the interrupt activation entries recorded in the intervening period. The policy of work delegation adopted by the system allows the periodic process to ensure the constant discharge of a guaranteed level of workload.

The correct implementation of this policy also requires assigning the periodic process a higher priority than those assigned to the sporadic processes, so that guaranteed work can be performed in preference to subsidiary activities.

The application is comprised by the following tasks and attributes. Static priorities are given based on the deadline monotonic scheduling [7], which the most optimal between the fixed priority algorithms [8].

| Task name | Task type | Period / Minimum inter-arrival time (ms) | Deadline (ms) | Priority |
|-----------------------|--------------------|--|---------------|----------|
| Regular_Producer | Cyclic | 1000 | 500 | 7 |
| On_Call_Producer | Sporadic | 5000 | 800 | 5 |
| Activation_Log_Reader | Sporadic | 3000 | 1000 | 3 |
| External_Event_Server | Interrupt sporadic | 5000 | 100 | 11 |

Table 1: Attributes of the tasks in the application [1]

Ada protected objects [RM 9.4] are used to ensure mutually exclusive access to shared resources, whereas protected entries are used only for task synchronization purposes where data exchange is involved.

In a real-time application, each protected object has a priority ceiling which represents the maximum priority of any task that calls the object. The Ada Real-Time Systems Annex supports the definition of Locking_Policy [RM D.3] and implements the Immediate Ceiling Priority Protocol (ICPP), usually called Priority Ceiling Protocol (PCP) in literature. It's one of the best Priority inheritance protocols, which allow a task to execute with an enhanced priority if it is blocking (or could block) a higher-priority task. To be specific, PCP reduces blocking to its minimum value: every job is blocked at most once for the duration of a critical section, no matter how many jobs conflict with it [9].

| Protected object names | User tasks | Ceiling priority |
|------------------------|---|---------------------------------|
| Request_Buffer | Regular_Producer (Deposit), On_Call_Producer (Extract) | 9 |
| Event_Queue | External interrupt (Signal), External_Event_Server (Wait) | System.Interrupt_Priority'First |
| Activation_Log | External_Event_Server (Write), Activation_Log_Reader (Read) | 13 |

Table 2: Attributes of the protected objects in the application [1]

1.2 The Board

We have preferred the usage of a bare-board STM32F429I-Discovery instead of the GNAT emulator because the measure time values had too much standard deviation.

It has ST-LINK debugger built-in.

2 System model and notation

The described application is a set of tasks executing in the same processor, grouped into entities called transactions [21]. Each transaction Γ_i is activated by a periodic sequence of external events with period T_i , and contains a set of tasks. Each task is released when a relative time offset elapses after the arrival of the external event. Each activation of a task releases the execution of one instance of that task, called a *job*.




Figure 2: Timeline of a system composed of transactions with offsets [20]

Figure 2 shows an example of such system: the horizontal axis represents time; down-pointing arrows represent the arrival of the external events associated to each transaction, while up-pointing arrows represent the activation times of each task; and shaded boxes represent task execution [20]. Each task has its own unique priority and in this example the task set is scheduled using a preemptive fixed priority scheduling.

Each task will be identified with two subscripts: the first one identifies the transaction to which it belongs, and the second one the position that the task occupies within the tasks of its transaction, when they are ordered by increasing offsets. In this way, τ_{ij} will be the j -th task of transaction Γ_i , with an offset of Φ_{ij} and a worst-case execution time of C_{ij} . In addition, we will allow each task to have jitter, that is to have its activation time delayed by an arbitrary amount of time between 0 and the maximum jitter for that task, which we will call J_{ij} . This means that the activation time of task τ_{ij} may occur at any time between $t_0 + \Phi_{ij}$ and $t_0 + \Phi_{ij} + J_{ij}$, where t_0 is the instant at which the external event arrived.

The reason for this is that tasks must execute in order, e.g. On_Call_Producer can start executing only after the preceding task Regular_Producer in the transaction has completed. The precedence constraints are modeled by assigning each task

an initial offset and a maximum jitter [21]. The initial offset Φ_{ij} of a periodic task is the instant of the first activation of the task. However, a task belonging to a transaction may start only after it has been activated and the preceding task in the transaction has completed execution. Hence maximum jitter is the maximum time interval it can occur from the task activation until the completion time of the preceding task in the transaction.

In addition to maximum jitter, tasks offsets are allowed to vary dynamically, from one activation to the next, within a minimum and a maximum value: $\Phi_{ij} \in [\Phi_{ij \min}, \Phi_{ij \max}]$. Dynamic offsets are useful in systems in which tasks suspend themselves, like in the case of protected object entries. The task `On_Call_Producer` τ_{i2} calls the protected entry `Extract` and suspends itself until the task `Regular_Producer` τ_{i1} replenishes the `Request_Buffer`. The activation time of `On_Call_Producer` depends on the completion time of the `Regular_Producer` and thus the offset for task τ_{i2} is variable in the interval $\Phi_{i2} \in [R_{i1 \min}, R_{i1 \max}]$, where $R_{i1 \min}$ and $R_{i1 \max}$ are respectively the best-case and worst-case response times of task `Regular_Producer`.

For each task τ_{ij} we define its response time as the difference between its completion time and the instant at which the associated external event arrived. The worst-case response time will be called R_{ij} . Each task has also an associated global deadline, D_{ij} , which is again relative to the arrival of the external event.

If tasks synchronize using shared resources in a mutually exclusive way, they will be using the aforementioned Priority Ceiling Protocol. The effects of lower priority tasks on a task under analysis τ_{ab} are bounded by an amount called the blocking term B_{ab} , calculated as the maximum of all the critical sections of lower priority tasks that have a priority ceiling higher than or equal to the priority of τ_{ab} .

2.1 Holistic analysis

Rate monotonic analysis (RMA) [7] allows an exact calculation of the worst-case response time of tasks in single-processor real time systems, including the effects of task synchronization, the presence of aperiodic tasks, the effects of deadlines before, at or after the periods of the tasks, precedence constraints and tasks with varying priorities, overhead analysis, etc. However, classic RMA [16] cannot provide exact solutions in systems in which tasks suspend themselves. Classic techniques for these systems are based on the assumption that all tasks are independent, and thus they lead to pessimistic results [20].

For building the worst-case scenario for a task τ_{ab} under analysis, the analysis must consider the critical instant that leads to the worst-case busy period. A task τ_{ab} busy period is an interval of time during which the CPU is busy processing task τ_{ab} or higher priority tasks. For tasks with offsets, it must take into account that the critical instant may not include the simultaneous activation of all higher priority tasks, as it was the case when all tasks were independent. The existence of offsets makes it impossible for some sets of tasks to simultaneously become active.

Works on such problem has been the base of holistic analysis, first proposed by Tindell and Clark [21] for distributed systems and later improved by Palencia and González [20] who called it Worst-Case analysis of Dynamic Offsets (WCDO). In such analysis, the worst-case response time of each task is used to set the offset and the jitter of the successive task in the same transaction. Then, the computation of worst-case response times is iterated until a stable solution is found. If response times are bounded, the holistic method is guaranteed to converge to a solution.

The MAST analysis tool implements both the latest offset-based WCDO techniques and the more pessimistic holistic approach [3], which is included in the toolset for completeness and comparison.

3 MAST

MAST [3] is a Modeling and Analysis Suite for Real-Time Applications and its main goal is to provide an open source set of tools that enables engineers developing real-time applications to check the timing behavior of their application, including schedulability analysis for checking hard timing requirements.

It is designed to handle both fixed priority and dynamic priority scheduled systems, although offset-based analysis for Earliest Deadline First scheduling is still missing as of the time of writing. However, within fixed priorities, different scheduling strategies are allowed, including preemptive and non preemptive scheduling, interrupt service routines, sporadic server scheduling, and periodic polling servers.

The MAST model is designed to handle both single-processor as well as multiprocessor or distributed systems. In both cases, emphasis is placed on describing event-driven systems in which each task may conditionally generate multiple events at its completion. A task may be activated by a conditional combination of one or more events. The external events arriving at the system can be of different kinds: periodic, unbounded aperiodic, sporadic, bursty, or singular (arriving only once).

The system model facilitates the independent description of overhead parameters such as processor overheads (including the overheads of the timing services). This frees us from the need to include all these overheads in the actual application model, thus simplifying it and eliminating a lot of redundancy.

MAST provides also a graphical editor to generate the system using the MAST ASCII description, but it's still immature to be reliable and the presence of several graphical bugs causes an annoying experience. A graphical display of results is also available.

3.1 The MAST Model

We now proceed to describe the MAST model of the example application. For a reference to the MAST syntax, visit "Description of the MAST Model" [4].

3.1.1 Processing Resources

Processing Resources represent resources that are capable of executing abstract activities, including conventional CPU processors. Among its attributes we have the range of priorities valid for normal operations on that processing resource, and the speed factor. We have left the default value as speed factor, meaning that execution times will be expressed as seconds.

Normally when dealing with hard real-time analysis, we would also define only the Worst-Case Execution Time (WCET) of the operations but, since we have dynamic offsets depending on them, we include both best and worst execution times because we don't know for sure that always having the WCET corresponds to worst system performance. We may for instance have anomalies as in the case of multiple processors [12]. Nevertheless, most of the time we had deterministic execution times with always the same number of CPU ticks or with a difference less than $1\mu s$, except for the Whetstone operations which showed more significant differences.

```
Processing_Resource (
  Type           => Regular_Processor ,
  Name           => cpu ,
  Max_Interrupt_Priority => 255 ,
  Min_Interrupt_Priority => 241 ,
  Worst_ISR_Switch   => 2.578E-06 ,
  System_Timer      =>
    ( Type           => Ticker ,
      Worst_Overhead => 3.844E-06 ,
      Period         => 0.001000 ) ,
  Speed_Factor     => 1.00 );
```

We have a board with only one CPU and interrupt ranges are taken from System package. Task priorities range from 1 to 240, whereas interrupt priorities go from 241 to 255. So we have at maximum 240 task priorities, if we ever would need more priorities we could use the technique described in [10]. The Interrupt Service Routine overhead is measured as the time taken to execute `Interrupt_Handler` in `System.BB.Board_Support` package, apart from the execution time of the user-defined interrupt handler. It takes into account the overhead of managing the Task and Interrupt Execution Clocks [19].

The system timer used by the board is Tick Scheduling [11], which is accounted in the analysis using the technique described in [17]. The period of the tick is 1ms, as defined in `System.BB.Board_Support` package, and the worst overhead is measured as the time taken to execute `Timer_Interrupt_Handler`, the trap handler defined in the same package for `Sys_Tick` trap.

```
Scheduler (
  Type           => Primary_Scheduler ,
  Name           => fps ,
  Host           => cpu ,
  Policy         =>
    ( Type           => Fixed_Priority ,
      Worst_Context_Switch => 3.090E-06 ,
      Max_Priority      => 240 ,
      Min_Priority      => 1 ));
```

We only have one primary scheduler in the system, no hierarchical scheduling. The context switch overhead is measured as time to trigger the context switch interrupt `Pend_SV` and the execution time of `Pend_SV_Handler` in `System.BB.CPU_Primitives.Context_Switch_Trigger` package, to save the active context and load the new one.

```

Scheduling_Server (
  Type           => Regular ,
  Name           => regular_producer ,
  Server_Sched_Parameters =>
    ( Type       => Fixed_Priority_Policy ,
      The_Priority => 7,
      Preassigned => YES),
  Scheduler      => fps );

Scheduling_Server (
  Type           => Regular ,
  Name           => on_call_producer ,
  Server_Sched_Parameters =>
    ( Type       => Fixed_Priority_Policy ,
      The_Priority => 5,
      Preassigned => YES),
  Scheduler      => fps );

Scheduling_Server (
  Type           => Regular ,
  Name           => activation_log_reader ,
  Server_Sched_Parameters =>
    ( Type       => Fixed_Priority_Policy ,
      The_Priority => 3,
      Preassigned => YES),
  Scheduler      => fps );

Scheduling_Server (
  Type           => Regular ,
  Name           => external_event_server ,
  Server_Sched_Parameters =>
    ( Type       => Fixed_Priority_Policy ,
      The_Priority => 11,
      Preassigned => YES),
  Scheduler      => fps );

Scheduling_Server (
  Type           => Regular ,
  Name           => interrupt_server ,
  Server_Sched_Parameters =>
    ( Type       => Interrupt_FP_Policy ,
      The_Priority => 241),
  Scheduler      => fps );

```

Each task is a Scheduling Server, whereas `interrupt_server` models the runtime which runs the Interrupt Service Routine using the technique described in [18].

```

Shared_Resource (
  Type       => Immediate_Ceiling_Resource ,
  Name       => request_buffer ,
  Ceiling    => 9,
  Preassigned => YES);

Shared_Resource (
  Type       => Immediate_Ceiling_Resource ,
  Name       => activation_log ,

```

```

Ceiling      => 13,
Preassigned => YES);

```

```

Shared_Resource (
  Type      => Immediate_Ceiling_Resource ,
  Name      => event_queue ,
  Ceiling   => 241,
  Preassigned => YES);

```

Protected objects are modeled as Shared Resources which use the Immediate Priority Ceiling Protocol, the same defined as Priority Ceiling Locking in Ada Reference Manual D.3.

```

Operation (
  Type      => Simple ,
  Name      => rb_deposit ,
  Worst_Case_Execution_Time => 2.000E-06,
  Shared_Resources_To_Lock  =>
    ( request_buffer ),
  Shared_Resources_To_Unlock =>
    ( request_buffer ));

```

```

Operation (
  Type      => Simple ,
  Name      => rb_extract ,
  Worst_Case_Execution_Time => 2.000E-06,
  Shared_Resources_To_Lock  =>
    ( request_buffer ),
  Shared_Resources_To_Unlock =>
    ( request_buffer ));

```

Protected methods are modeled as simple operations which lock and unlock the protected object resource. The execution time is measured from the first line of the method to the last one, so it doesn't include the runtime overhead associated with invoking protected methods.

```

Operation (
  Type      => Enclosing ,
  Name      => ocp_start ,
  Worst_Case_Execution_Time=> 6.000E-06,
  Composite_Operation_List =>
    ( rb_deposit ));

```

```

Operation (
  Type      => Enclosing ,
  Name      => rb_extract_enclosing ,
  Worst_Case_Execution_Time=> 7.000E-06,
  Composite_Operation_List =>
    ( rb_extract ));

```

For each protected method there is an Enclosing operation which takes into account the overhead associated with invoking protected methods. Sometimes it's already a method defined by GEE, other times it's defined in the model on purpose. By doing so we can define the more complex methods as Composite operations, which have the execution time as the sum of the execution times of the comprised operations.

```

Operation (
  Type      => Simple ,
  Name      => rp_small_whetstone ,
  Worst_Case_Execution_Time => 0.019363);

```

```

Operation (
  Type      => Composite ,
  Name      => rp_operation ,

```

```

Composite_Operation_List =>
( rp_small_whetstone ,
  due_activation ,
  ocp_start ,
  check_due ,
  alr_signal ,
  put_line ));

Operation (
  Type                => Composite ,
  Name                => regular_producer ,
  Composite_Operation_List =>
    ( overrun_detection ,
      rp_operation ,
      delay_until ));

```

Therefore, by changing the workload parameter of Small_Whetstone in the application implementation, we will be able to test different utilisation of the system with likewise ease in updating the MAST model. The Whetstone execution time is proportional to the workload parameter. If we wanted to try what happens by increasing the load of factor 10, we would just increase the WCET to 0.19363, without the need to measure again all the Enclosing operations, since all the methods which use the Whetstone are defined as Composite. However we have been careful to avoid forgetting any overhead in a Composite method and make sure they are not impacted by any change of the Whetstone workload. For instance, if we had defined regular_producer as Enclosing we could have defined as composed of only rp_operation and then measure its WCET, which will implicitly count also overrun detection and delay queue overhead. By instead defining it as Composite, we have been careful to define the simple operations overrun_detection and delay_until to include their execution time, using the technique described in [17].

```

Transaction (
  Type                => regular ,
  Name                => rp_transaction ,
  External_Events =>
    ( ( Type          => Periodic ,
        Name          => e1 ,
        Period        => 1.000 ,
        Max_Jitter    => 0.000 ,
        Phase         => 0.000)),
  Internal_Events =>
    ( ( Type => Regular ,
        Name => rpol ,
        Timing_Requirements =>
          ( Type          => Hard_Global_Deadline ,
            Deadline      => 0.500000 ,
            Referenced_Event => e1))),
  Event_Handlers =>
    ( (Type          => System_Timed_Activity ,
        Input_Event  => e1 ,
        Output_Event => rpol ,
        Activity_Operation => regular_producer ,
        Activity_Server  => regular_producer )));

```

The main event stream is modeled as a transaction activated by a periodic external event, with period of 1s. The event is handled by the regular_producer operation by the task of the same name. The Event Handler is of type System_Timed_Activity to take into account the jitter and the overhead caused by the tick scheduling

```

Transaction (
  Type                => regular ,
  Name                => ocp_transaction ,
  External_Events =>
    ( ( Type          => Sporadic ,
        Name          => ocp_activation ,

```



```

    Avg_Interarrival => 5.000,
    Distribution      => UNIFORM,
    Min_Interarrival => 5.000)),
Internal_Events =>
  ( ( Type => Regular,
    Name => ocpol,
    Timing_Requirements =>
      ( Type => Hard_Global_Deadline,
        Deadline => 0.800000,
        Referenced_Event => ocp_activation))),
Event_Handlers =>
  ( (Type => Activity,
    Input_Event => ocp_activation,
    Output_Event => ocpol,
    Activity_Operation => on_call_producer,
    Activity_Server => on_call_producer)));

```

The sporadic On Call Producer event stream is modeled as activated by a bounded aperiodic event, with minimum interarrival time of 5s and uniform distribution. Actually we know that the interarrival time is precisely 5s, thus the same value as average interarrival. Similar modeling has been done for the Activation Log Reader sporadic task.

```

Transaction (
  Type      => regular,
  Name      => event_queue_interrupt,
  External_Events =>
    ( ( Type => Sporadic,
      Name => button_click,
      Avg_Interarrival => 0.000,
      Distribution      => UNIFORM,
      Min_Interarrival => 5.000)),
  Internal_Events =>
    ( ( Type => Regular,
      Name => eqol),
      ( Type => Regular,
        Name => eqo2,
        Timing_Requirements =>
          ( Type => Hard_Global_Deadline,
            Deadline => 0.100000,
            Referenced_Event => button_click))),
  Event_Handlers =>
    ( (Type => Activity,
      Input_Event => button_click,
      Output_Event => eqol,
      Activity_Operation => eq_signal,
      Activity_Server => interrupt_server),
      (Type => Activity,
        Input_Event => eqol,
        Output_Event => eqo2,
        Activity_Operation => external_event_server,
        Activity_Server => external_event_server)));

```

The blue button interrupt event stream is modeled as a triggered by a sporadic event of 5s as minimum interarrival time and it's first handled by the `interrupt_server` which runs the ISR at interrupt priority level and then by the user-defined `external_event_server` at task priority-level.

4 Overrun detection

— *Overrun.ads*

```

with Ada.Real_Time;
with Ada.Execution_Time;

package Overrun is
  type Limits_Array is array (0 .. 2) of Ada.Execution_Time.CPU_Time;

  procedure Start (Index : Natural; Budget : Ada.Real_Time.Time_Span);
  procedure Check (Index : Natural);
end Overrun;

```

— *Overrun.adb*

```

with Ada.Real_Time; use Ada.Real_Time;
with Ada.Execution_Time; use Ada.Execution_Time;

package body Overrun is
  use Ada.Real_Time;
  use Ada.Execution_Time;

  Limits : Limits_Array := (CPU_Time_First, CPU_Time_First, CPU_Time_First);

  procedure Start (Index : Natural; Budget : Time_Span) is
  begin
    Limits (Index) := Ada.Execution_Time.Clock + Budget;
  end Start;

  procedure Check (Index : Natural) is
  begin
    if Ada.Execution_Time.Clock > Limits (Index) then
      raise Program_Error with "Detected_overflow";
    end if;
  end Check;
end Overrun;

```

Inspiration from [6]. The measured execution times include overrun detection overhead for Regular Producer, On Call Producer and Activation Log Reader.

5 Execution times

Execution times are measured using two packages `System_Overhead` and `Task_Overhead` using Semihosting.

6 MAST analysis

As of the time of writing, MAST is at version 1.5.1 and supports the following analysis tools.

The transaction which defines the button interrupt event stream is considered as a linear transaction, which only has one external event and that its Event handlers are all Activities, but unfortunately it is not a simple transaction, a continuous sequence of activities executed by the same server. We have one server which handles the ISR and another one the associated task of External Event Server.

This means we cannot use the classic Rate Monotonic algorithm [7], only offset-based (cite ?) and holistic analysis [13]. That's probably because Rate Monotonic assumes independent tasks, but our interrupt transaction is composed by an ISR in a `Interrupt_FP_Server`, followed by the interrupt handler in `External Event Server` which is activated by the former and thus it depends on it for its activation. We decided to stick to only holistic analysis because it supports both FPS and EDF, whereas offset-based fallbacks to holistic analysis with EDF processing resources [5]. Nevertheless it doesn't make much difference which one is used between holistic and offset-based since we run the system on a single processor, not on a distributed system (?).

Table 1. Fixed-priority schedulability analysis tools

| Technique | Single-Processor | Multi-Processor | Simple Transact. | Linear Transact. | Multipath Transact. |
|------------------------|------------------|-----------------|------------------|------------------|---------------------|
| Classic Rate Monotonic | ✓ | | ✓ | | |
| Varying Priorities | ✓ | | ✓ | ✓ | |
| Holistic | ✓ | ✓ | ✓ | ✓ | ✓ |
| Offset Based | ✓ | ✓ | ✓ | ✓ | |

Table 2. EDF schedulability analysis tools

| Technique | Single-Processor | Multi-Processor | Simple Transact. | Linear Transact. | Multipath Transact. |
|-----------------------|------------------|-----------------|------------------|------------------|---------------------|
| Single Processor | ✓ | | ✓ | | |
| EDF_Within_Priorities | ✓ | | ✓ | | |
| Holistic_Local | ✓ | ✓ | ✓ | ✓ | ✓ |
| Holistic_Global | ✓ | ✓ | ✓ | ✓ | ✓ |
| Offset Based | ✓ | ✓ | ✓ | ✓ | |

Figure 3: MAST analysis tools [4].

| Transaction | Worst case response time (s) | Slack | Worst blocking time (s) |
|-----------------------|------------------------------|----------|-------------------------|
| rp_transaction | 0.020393 | 2477.0% | 2.000E-06 |
| ocp_transaction | 0.026525 | 10852.0% | 1.000E-06 |
| alr_transaction | 0.030109 | 27088.7% | 0.00 |
| event_queue_interrupt | 3.818E-05 | N/A | 1.000E-06 |

Table 3: Holistic analysis results for FPS

7 FPS analysis

The system slack is 2401.2% and total utilisation 2.59%.

We then increase the Whetstone workload of factor 24 in the first three transactions since 2477.0% is the smallest slack of three transactions. We leave the Event Queue interrupt unchanged. The new results are as follows:

| Transaction | Worst case response time (s) | Slack | Worst blocking time (s) |
|-----------------------|------------------------------|-------------|-------------------------|
| rp_transaction | 0.487017 | 2.34% | 2.000E-06 |
| ocp_transaction | 0.664777 | 75.39% | 1.000E-06 |
| alr_transaction | 0.754207 | 273.44% | 0.00 |
| event_queue_interrupt | 3.818E-05 | >=100000.0% | 1.000E-06 |

Table 4: Holistic analysis results for FPS

Blocking times have not changed because protected operations are same as before.

The system slack is 2.80% and total utilisation 55.33%. The theoretical CPU utilisation upper bound [14] is 0.779 for tasks with same deadline as the period, but we have to use the technique shown in [15]. We now increase the workloads of factor 25 instead of 24%, assuming we had a slack of 2500%.

The system slack is -1.16% and total utilisation 57.53%, which exceeds the theoretical limit (?). This means that the actual execution should also overrun the deadline and it is indeed what happened on our board. The first job of Regular Producer raised the overrun detection Program Error.

| Transaction | Worst case response time (s) | Slack | Worst blocking time (s) |
|-----------------------|------------------------------|----------|-------------------------|
| rp_transaction | 0.506453 | -1.56% | 2.000E-06 |
| ocp_transaction | 0.691366 | -100.00% | 1.000E-06 |
| alr_transaction | 0.784369 | -100.00% | 0.00 |
| event_queue_interrupt | 3.818E-05 | -100.00% | 1.000E-06 |

Table 5: Holistic analysis results for FPS

The protected methods execution times are so small compared to Whetstone that even if we model it badly, it doesn't matter. But actually we are being too pessimistic because our application has dependency chains, tasks don't compete for the same resource. They synchronize. Let's try putting Whetstone inside the protected methods, now too pessimistic analysis matters. In particular there are 500ms after Regular Producer completion and its next release. If On Call Producer and Activation Log Reader spend less than 500ms together in protected methods, we know for sure that Regular Producer is never blocked by them. Pessimist analysis may consider the blocking time however. Maybe we can model protected methods as message communication overhead and use holistic analysis for distributed systems. We may also increase execution times of On Call Producer and Activation Log Reader because maybe the analysis is too pessimistic with the critical instants.

Two pessimistic points:

1. Blocking time
2. Critical instant/busy period

References

- [1] A Burns, B Dobbing, T Vardanega. Guide for the use of the Ada Ravenscar Profile in high integrity systems. In *University of York Technical Report YCS-2003-348*. January 2003.
- [2] ARM KEIL. What is semihosting? In *Compiler User Guide*. http://www.keil.com/support/man/docs/armcc/armcc_pge1358787046598.htm.
- [3] M. GonzAlez Harbour, J.J. GutiCrrez Garcia, J.C. Palencia GutiCrrez, and J.M. Drake Moyano. MAST Modeling and Analysis Suite for Real Time Applications. In *Proceedings 13th Euromicro Conference on Real-Time Systems*. 2001.
- [4] J. M. Drake, M. G. Harbour, J. J. Gutiérrez, P. L. Martínez, J. L. Medina, J. C. Palencia Description of the MAST Model. https://mast.unican.es/mast_description.pdf
- [5] J. M. Drake, M. G. Harbour, J. J. Gutiérrez, P. L. Martínez, J. L. Medina, J. C. Palencia MAST README. <https://mast.unican.es/README.txt>
- [6] Juan Zamorano, Alejandro Alonso, José Antonio Pulido, Juan Antonio de la Puente. Implementing Execution-Time Clocks for the Ada Ravenscar Profile. In *Reliable Software Technologies - Ada-Europe 2004*. pp 132-143. Ada-Europe 2004.
- [7] Jane W. S. W. Liu. Rate-Monotonic and Deadline-Monotonic Algorithms. In *Real-Time Systems*. pp 118-119. 2001.
- [8] Jane W. S. W. Liu. Optimality of the RM and DM algorithms. In *Real-Time Systems*. pp 118-119. 2001.
- [9] Jane W. S. W. Liu. Basic Priority Ceiling Protocol - Duration of Blocking. In *Real-Time Systems*. pp 295-296. 2001.
- [10] Jane W. S. W. Liu. Limited-Priority Levels. In *Real-Time Systems*. pp 166-168. 2001.
- [11] Jane W. S. W. Liu. Tick Scheduling. In *Real-Time Systems*. pp 168-171. 2001.
- [12] Jane W. S. W. Liu. Anomalous Behavior of Priority-Driven Systems. In *Real-Time Systems*. pp 72-73. 2001.
- [13] KenTindell, JohnClark. Holistic schedulability analysis for distributed hard real-time systems. In *Microprocessing and Microprogramming*. Volume 40, Issues 2-3, pp 117-134. April 1994.
- [14] C. L. Liu, James W. Layland. Scheduling Algorithms for Multiprogramming in a Hard-Real-Time Environment. In *Journal of the ACM*. Volume 20 Issue 1, pp 46-61. Jan. 1973.
- [15] Klein, M., Ralya, Th., Pollak, B., Obenza, R., Harbour, M.G. . Using Utilization Bounds for Each Event when Deadlines Are Within the Period. In *A Practitioner's Handbook for Real-Time Analysis* . chapter 4.1.2. 1993.
- [16] Klein, M., Ralya, Th., Pollak, B., Obenza, R., Harbour, M.G. . Designing Tasks that Must Synchronize to Share Common Data. In *A Practitioner's Handbook for Real-Time Analysis* . chapter 5.2. 1993.

- [17] Klein, M., Ralya, Th., Pollak, B., Obenza, R., Harbour, M.G. . Effects of Operating System and Runtime Services on Timing Analysis. In *A Practitioner's Handbook for Real-Time Analysis* . chapter 7. 1993.
- [18] Klein, M., Ralya, Th., Pollak, B., Obenza, R., Harbour, M.G. . Use an Interrupt Handler. In *A Practitioner's Handbook for Real-Time Analysis* . chapter 5.3.5.1. 1993.
- [19] Kristoffer Nyborg Gregertsen, Amund Skavhaug. Implementation and Usage of the new Ada 2012 Execution Time Control Features. In *Ada User Journal*. 2011.
- [20] J.C. Palencia ; M. Gonzalez Harbour. Schedulability analysis for tasks with static and dynamic offsets. In *Proceedings 19th IEEE Real-Time Systems Symposium*. 1998.
- [21] K. Tindell. Adding Time - Offsets to Schedulability Analysis. Technical Report YCS 221, Dept. of Computer Science, University of York, England, January 1994.